

Matthew Wisdom

matthewwisdom11@gmail.com | +234 912 897 6463 | matt-wisdom.github.io | github.com/matt-wisdom

Software and AI engineer with 5 years of experience specializing in Full-Stack Web Development, LLM orchestration, production ML systems, and high-throughput APIs. Proven track record of reducing LLM costs by 50%+, cutting API latency 3×, and shipping AI pipelines that serve thousands of daily requests.

TECHNICAL SKILLS

Languages: Go, Python, JavaScript/TypeScript, SQL, Bash

AI & Web: LangGraph, LangChain, RAG (Qdrant), vLLM, Ray, MLFlow, PyTorch, FastAPI, Go Fiber, RabbitMQ, Redis, PostgreSQL, NodeJS, MCP, VueJS, NuxtJS

Cloud & Ops: AWS (EC2, Lambda, ECS), Docker, Terraform, Prometheus, Grafana, LangFuse, OpenTelemetry, GitHub Actions

EXPERIENCE

Backend AI Engineer · KwevoxAI (Remote, Nigeria) Jan 2025 – Present

- Architected a multi-stage AI generation pipeline using LangGraph for LaunchpadNG, coordinating LLMs to autonomously plan, generate code, and self-correct website and e-commerce deployments — cutting end-to-end site generation time by ~70% to under 5 minutes; generated sites now serve thousands of weekly requests from business customers.
- Engineered production RAG systems using Qdrant vector DB with intelligent rate limiting, response caching, and smart model routing — reducing LLM inference costs by 50% while maintaining retrieval quality at scale.
- Reduced API latency ~3× on high-traffic endpoints through Redis caching and PostgreSQL schema refactoring.
- Deployed a monitoring infrastructure with Grafana, Cloudwatch and Langfuse for granular AI telemetry, actively tracking daily events and maintaining **99.95%** uptime.

Fullstack & AI Engineer · Home Team Management LLC · Remote (US) Oct 2023 – Jan 2025

- Built a Go-based concurrent auto-fill API delivering predictive field suggestions at sub-50ms latency for military-scale administrative forms.
- Developed HeroForms.io, digitizing military workflows with encrypted e-signing and automated high-throughput document routing.
- Built and maintained DIYPayroll, a payroll and tax compliance system ensuring zero-error processing for \$50,000+ USD in monthly transactions.
- Designed and deployed an Isolation Forest anomaly detection model for payroll transaction monitoring — configured for high-recall compliance use (0.93 recall, 0.30 precision), flagging 4+ critical anomalies in the first month of operation and preventing downstream compliance breaches.

Backend Engineer (Contract) · US Startup Clients via Turing.com · Remote Feb 2023 – Apr 2024

- Delivered AI document processing pipelines and REST APIs serving 15,000+ daily requests with 99.9% uptime.
- Managed LLM serving infrastructure using Ray and vLLM; improved query latency 15% via PostgreSQL query optimization.

Freelance Backend Engineer · BetterCNC · Remote Jan 2023 – Apr 2023

- Built ETL pipelines using DuckDB and VueJS to ingest and visualize CNC machine job data, replacing manual tracking with real-time operational analytics.

Freelance Backend Engineer · [Lions PromoBot](#) · Remote

Apr 2022 – Jan 2023

- Scaled an asynchronous Python backend to handle 5,000+ concurrent requests, facilitating a 100× increase in user load while keeping query latency under 100ms.

Software Support Engineer · [Fixexception.com](#) · (Remote, Ukraine) Jan 2021 - March 2022

- Troubleshoot and resolved code-level exceptions for a global user base across Python and JavaScript.
- Debugged user-submitted scripts to identify root causes of logical and syntax errors, providing optimized code solutions.
- Communicated debugging strategies and technical fixes clearly to developers of varying skill levels.

EDUCATION

B.Eng, Mechatronics and Systems Engineering

Abubakar Tafawa Balewa University, Nigeria